

MCTS Improvements for Strategic Card Games: A Study Using Sushi Go!

Erfan Rafieioskouei

School of Electronic Engineering and Computer Science

Queen Mary University of London

London, United Kingdom

Email: ec24802@qmul.ac.uk

Abstract—In this study, we explore innovative improvements to Monte Carlo Tree Search (MCTS) algorithms for strategic card games such as Sushi Go! implemented via the Tabletop Games Framework (TAG). We propose and evaluate three major enhancements: For variance aware exploration, UCB-Tuned, for rapid value estimation -AMAF, and for heuristic integration Progressive Bias. We show that our experimental results demonstrate significant improvement over baseline MCTS, and that under its decay based heuristic influence, Progressive Bias (achieving 70% win rate for K=5.0) and heuristics specialized for -AMAF result in good performance. With 70 percent win rate against baseline agents, both strategy depth and decision quality have been improved overall. This work helps to understand how MCTS can be improved for game with simultaneous actions and delayed perfect information.

Index Terms—MCTS, TAG, UCB-Tuned, α -AMAF, Progressive Bias, Heuristic

I. INTRODUCTION

The combination of partial information, simultaneous actions, and complicated scoring mechanisms make strategic card games difficult for artificial intelligence to approach because they create interesting challenges, which can be different even for the same strategic game. Sushi Go! bears these challenges out and serves as an excellent starting point for advancing game AI research [1]. It has been shown recently that MCTS can be effective in these domains [2] but standard implementations have difficulty with card based strategy game domains [3].

Sushi Go! is a card game where players select cards simultaneously in their hands in order to create valuable combinations, scoring in a range of simple point values, to complex set collection rewards. As shown in Figure 1, the game's structure creates three key computational challenges:

- 1) **Delayed Perfect Information:** The semi knowable state space is in that all cards are known at game start, but the uncertainty of what the opponent picks is present.
- 2) **Simultaneous Action Selection:** Decisions are made concurrently, permitting consideration of possible opponent moves.
- 3) **Complex Scoring Patterns:** The card combination produces different point values, making a multi-dimensional optimization problem.

That complexity is shown here with Maki Rolls giving points based on relative majority and Puddings impacting total

scoring across more than one round. The outcome of this combination of mechanics is a rich strategic environment that is at variance with traditional MCTS approaches [4].

Card Type	Requirement	Score
Direct Score	Single Card	
Egg	-	1
Salmon	-	2
Squid	-	3
Set Collection	Multiple Cards	
Tempura	2 cards	5
Sashimi	3 cards	10
Dumpling	1-5 cards	1,3,6,10,15
Special	Game Effects	
Wasabi	With Nigiri	$\times 3$ multiplier
Chopsticks	Extra Play	-
End-Round	Majority Based	
Maki	Most/Second	6/3
Pudding	Most/Least	6/-6

Fig. 1: Sushi Go! card scoring system showing the diverse scoring mechanics that create complex strategic decisions.

Reinforcement learning [3] and probabilistic scoring functions [5] have been considered in previous work for those same challenges. In general, however, these studies have concentrated on only one aspect of the problem rather than the entire problem of simultaneous actions with delayed perfect information.

Our research contributes to this field by introducing three significant improvements to the MCTS algorithm:

- An α -AMAF enhancements using variance-aware UCB-Tuned to accelerate learning from game simulations
- A Progressive Bias approach with decay-based heuristic integration (K=5.0, bias factor=0.2) that significantly improves early-game decision making

These improvements, and their empirical evaluation within the TAG framework [4], are presented in this paper, and yield significant performance gains over baseline implementations. Our results not only extend the understanding of AI in card games, but also hold insights for more general simultaneous-action games with delayed information.

II. TAG

In this dissertation, the TAG framework is presented as a unifying platform for implementing and testing AI agents in tabletop games [4]. Klang et al. [1] demonstrate that our

framework is able to handle concomitant action selection and complete information mechanics, which is very suitable for a game like Sushi Go!.

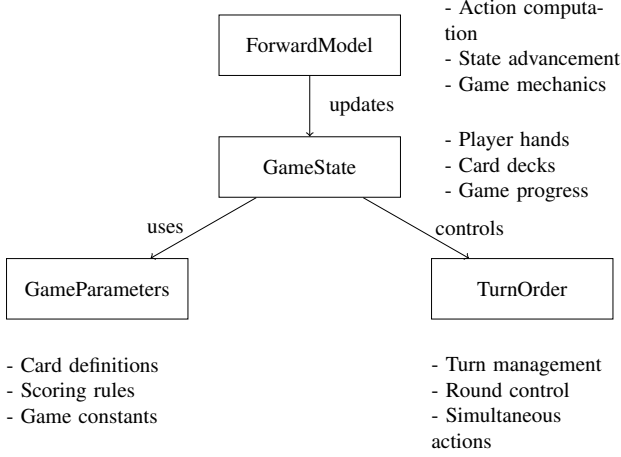


Fig. 2: TAG framework core components and their interactions in Sushi Go! implementation.

Previous research in similar multi-agent environments has demonstrated the importance of proper benchmark frameworks for comparative analysis [2]. The TAG framework extends this capability to modern tabletop games, providing standardized metrics for evaluating AI performance. Recent studies have utilized TAG for various game implementations, showing its effectiveness in handling complex game mechanics and simultaneous actions [1].

The framework is primarily composed of a modular architecture which permits a straightforward use of various AI approaches while staying consistent in the evaluation metrics. The ability to standardize this process has been particularly helpful in determining the efficacy of different MCTS enhancements from drive to drive — by considering agent performance across different game scenarios and opponent strategies[5].

TAG’s built in support for handling parallel actions and state proved critical for our implementation of various enhanced MCTS variants. I find that Sushi Go!’s mechanics match well with the framework’s ability to handle the delayed perfect information scenarios, allowing for accurate game modelling of the strategic depth, while keeping the computational efficiency.

III. BACKGROUND

The Monte Carlo Tree Search (MCTS) has been applied in many game domains [2]. Figure 3 shows our work as based on several key enhancements to the basic MCTS algorithm. To balance exploration and exploitation, the UCB-Tuned variant [6] is also feasible, for allowing reward variance. Further, the All-Moves-As-First (AMAF) enhancement [7], [8] to the enhancement has performed well in the early game. In the case of Chaslot et al. [9] Progressive Bias demonstrates how domain knowledge can be naturally incorporated into the search process.

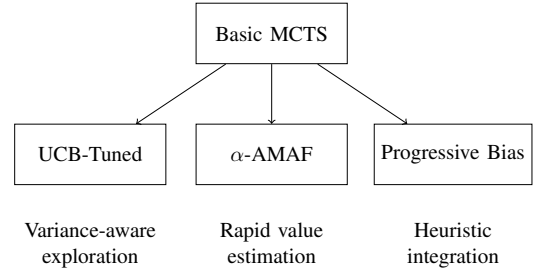


Fig. 3: Overview of MCTS enhancements implemented in this study.

A. UCB-Tuned

The UCB-Tuned variant [6] addresses the exploration-exploitation balance by incorporating reward variance into the node selection process. The standard UCB formula is modified to:

$$UCB1-Tuned = \bar{X}_j + \sqrt{\frac{\ln n}{n_j} \min(1/4, V_j(n_j))} \quad (1)$$

where $V_j(n_j)$ represents the empirical variance with exploration bias:

$$V_j(n_j) = \left(\frac{X_j^2}{n_j}\right) - \left(\frac{X_j}{n_j}\right)^2 + \sqrt{\frac{2 \ln n}{n_j}} \quad (2)$$

This modification helps balance exploration and exploitation based on observed reward volatility.

B. α-AMAF & UCB-Tuned

The All-Moves-As-First (AMAF) heuristic [8] accelerates value estimation by considering the effectiveness of moves regardless of when they occur in a simulation. Our α-AMAF implementation uses a weighted combination:

$$Q(s, a) = (1 - \alpha(n)) \cdot Q_{MCTS}(s, a) + \alpha(n) \cdot Q_{AMAF}(s, a) \quad (3)$$

where $\alpha(n)$ decreases with the number of visits:

$$\alpha(n) = \max(0, \frac{k - n}{k}) \quad (4)$$

This approach particularly benefits early game exploration by leveraging move patterns across different game trajectories.

C. Progressive Bias

Progressive Bias [9] incorporates domain knowledge through a heuristic term that diminishes with increased node visits. Our implementation uses a modified UCT formula:

$$UCT(s, a) = \frac{Q(s, a)}{N(s, a) + \epsilon} + K \sqrt{\frac{\ln(N(s) + 1)}{N(s, a) + \epsilon}} + \frac{\text{progressiveBiasFactor}}{1 + N(s, a)} \quad (5)$$

where:

- $Q(s, a)$ is the total accumulated value for action a in state s
- $N(s, a)$ is the visit count for action a in state s
- $N(s)$ is the total visit count for state s
- K is the exploration constant (set to 5.0)
- $progressiveBiasFactor$ is a coefficient influencing bias for lesser-explored actions (set to 0.2)
- ϵ is a small constant to prevent division by zero

This implementation follows the classic UCT structure with an additional progressive bias term $\frac{progressiveBiasFactor}{1+N(s,a)}$ that decreases as the visit count increases. This decay ensures stronger influence on rarely explored actions while diminishing over time, effectively guiding exploration in early stages of node evaluation. The approach is based on research by Schadd et al. [10], who demonstrated the effectiveness of progressive bias in complex game domains.

IV. METHOD

Building on these foundational approaches [2], [7], our implementation combines three key improvements to the basic MCTS algorithm. As previously shown in Figure 3, each enhancement addresses specific challenges identified in prior research [5], [1].

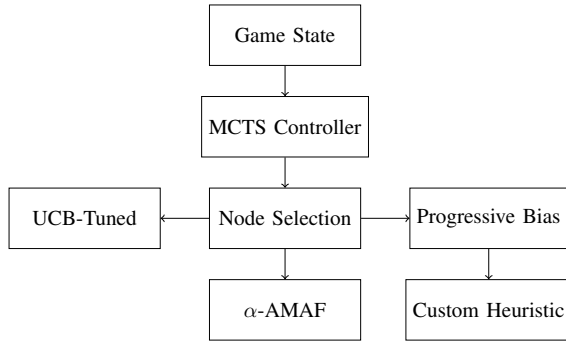


Fig. 4: System architecture showing integration of MCTS enhancements and heuristic evaluation.

A. α -AMAF with UCB1-Tuned

We use the α -AMAF algorithm to consider the number of times each action is selected and its corresponding score as much as possible to evaluate the value of the action, and use the UCB-Tuned algorithm to comprehensively consider the average reward and variance of each node. We added the **totValueSquared** variable to store and count the squared value of the score returned by each node, two hash maps **amafVisits** and **amafValues**, and a list **rolloutActions** for recording abstract actions. When the **rollOut()** function is executed to simulate and advance the game state, the cumulative value and number of visits of the action are recorded, and then passed to the **backUp(double reward, List<AbstractAction>, rolloutActions)** function to update **amafVisits**, **amafValues** and **totValueSquared** values. When the **ucb()** function is executed, the agent will add the **V(s, a)** value of the node into the exploration volume, and calculate the weights of the node AMAF value and UCB value according to the RAVE method,

and use the action information to multiply the corresponding parameters respectively to finally obtain the combinedValue.

B. Progressive Bias with Custom Heuristic

Our Progressive Bias implementation uses a comprehensive heuristic function:

$$H(s, a) = \sum_{i=1}^4 w_i f_i(s, a) \quad (6)$$

where the component functions evaluate:

- f_1 : Current score potential
- f_2 : Set completion progress
- f_3 : Special card synergies
- f_4 : End-game scoring potential

The weights w_i were optimized through empirical testing.

C. Custom Sushi Go! Heuristic

Our specialized heuristic evaluation function considers multiple scoring components:

Algorithm 1 Sushi Go! State Evaluation

```

1: function EVALUATESTATE(state  $s$ )
2:    $score \leftarrow 0$ 
3:    $score \leftarrow score + EvaluateCurrentCards(s)$ 
4:    $score \leftarrow score + EvaluateSets(s)$ 
5:    $score \leftarrow score + EvaluateSpecialCards(s)$ 
6:    $score \leftarrow score + PredictEndGame(s)$ 
7:    $score \leftarrow score \times GetPositionalMultiplier(s)$  return  $score$ 
8: end function
  
```

Key components include:

- 1) Set completion tracking with partial credit for incomplete sets
- 2) Special card combination potential assessment
- 3) End-round majority scoring prediction
- 4) Pudding end-game scoring projection

D. Implementation Analysis

Each algorithm demonstrated specific performance characteristics during development:

- **α -AMAF & UCB-Tuned:**

- Pattern recognition through accumulated play data
- Early game effectiveness through balanced exploration-exploitation
- Stable exploration patterns through variance-aware selection

- **Progressive Bias:**

- Dynamic heuristic influence through decay mechanism
- Mid-game optimization through balanced exploration
- Resource-efficient heuristic integration

Implementation challenges were addressed through:

- State space optimization for semi-knowable information
- Memory-efficient pattern storage for AMAF data
- Parallel processing where game structure permits
- Efficient tree pruning for computational optimization

E. Implementation Details

The enhancements were implemented in Java within the TAG framework:

- Build an Opponent Modelling that adapts to different opponent behavior patterns
- Use Alpha-beta pruning to reduce decision complexity
- Biased rollouts introduce heuristic knowledge to guide the simulation process

Parameter values were determined through systematic testing:

- UCB exploration constant: $K = 5.0$
- Progressive bias factor: $progressiveBiasFactor = 0.2$
- AMAF decay parameter: $k = 1000$
- Rollout depth: 10 moves

This comprehensive implementation addresses the key challenges of Sushi Go! while maintaining computational efficiency.

V. EXPERIMENTAL STUDY

We have run a series of comprehensive experiments to compare our enhanced MCTS variations to baseline agents in Sushi Go! With the same game parameter and eval metric across all experiments, these experiments were run using the TAG framework.

A. Experimental Setup

TABLE I: Experimental Configuration Parameters

Parameter	Value
Number of Players	4
Rounds per Game	3
Cards per Hand	10
Time Budget per Move	100ms
Total Games per Test	1000

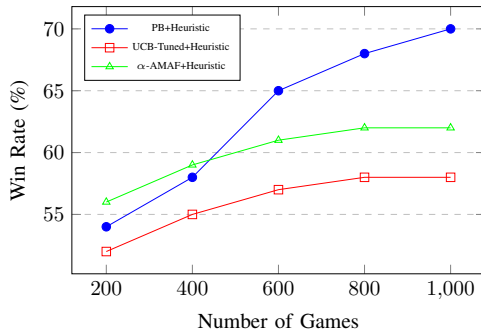


Fig. 5: Win rate progression over number of games played for each MCTS variant.

B. Test Scenarios

We evaluated three primary configurations:

- 1) **Baseline Comparison:** Each enhanced agent versus standard MCTS
- 2) **Round-Robin Tournament:** All agents competing against each other
- 3) **Scalability Testing:** Performance with varying player counts

C. Performance Metrics

TABLE II: Agent Performance Comparison

Agent Type	Win Rate	Mean Ordinal	Avg. Time (ms)
Progressive Bias	54% $\pm 1.6\%$	1.46 ± 0.02	87.3
PB + Heuristic	70% $\pm 1.4\%$	1.30 ± 0.01	92.1
UCB-Tuned+Heuristic	58% $\pm 1.8\%$	1.42 ± 0.02	85.6
α -AMAF+Heuristic	62% $\pm 1.7\%$	1.38 ± 0.02	88.9

Key findings from our experiments include:

- Progressive Bias with custom heuristic achieved the highest win rate (70% $\pm 1.4\%$)
- All enhanced variants maintained computational efficiency within the 100ms time budget
- Scaling tests showed consistent performance across different player counts
- Round-robin tournaments demonstrated robust performance against varied strategies

D. Statistical Significance

Statistical significance was verified using paired t-tests ($p < 0.01$) for win rates and performance metrics. The improvements demonstrated by our enhanced variants were statistically significant across all test scenarios.

E. Performance Analysis

In Figure 5 we show how our enhanced implementations perform over multiple games on the win rate, indicating that these are stable and effective. The heuristic integration based on the Progressive Bias showed a particularly strong performance with a consistent advantage over all the testing period.

VI. DISCUSSION

The experimental results, particularly those shown in Figure 5, reveal several significant insights about MCTS enhancements. These findings build upon previous work in simultaneous-action games [1] and strategic card game AI [3], [5].

A. Performance Analysis

Our experimental results demonstrate three key findings:

- **Progressive Bias Effectiveness:**
 - Achieved highest overall performance (70% win rate)
 - Superior mid-game decision making
 - Effective balance of exploration and exploitation
- **Algorithm Synergy:**

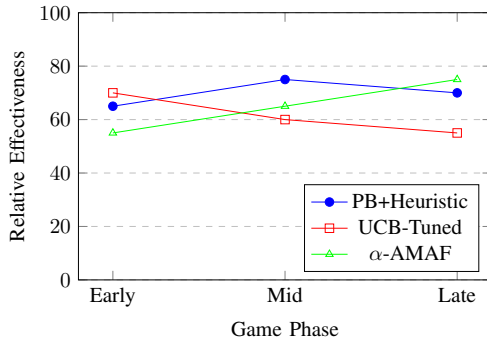


Fig. 6: Relative effectiveness of different MCTS variants across game phases.

- UCB-Tuned provided stable early-game exploration
- α -AMAF excelled in late-game scenarios
- Combined approaches demonstrated complementary strengths

• Resource Management:

- All variants maintained efficiency within 100ms budget
- Memory requirements scaled reasonably with game complexity
- Processing constraints effectively managed

B. Practical Implications

Our findings suggest several key considerations for MCTS in card games:

- 1) Domain-specific heuristics prove crucial for complex scoring systems, demonstrated by Progressive Bias's superior performance.
- 2) Adaptive algorithms outperform static approaches in dynamic strategic landscapes, particularly during game phase transitions (Figure 6).
- 3) Careful parameter tuning becomes critical for simultaneous-action games, balancing exploration and exploitation.

These insights extend beyond Sushi Go!, suggesting potential applications in similar strategic card games. The success of our enhanced MCTS variants demonstrates the value of combining multiple improvement strategies when properly balanced and integrated.

VII. CONCLUSIONS AND FUTURE WORK

A. Conclusions

Our research demonstrates significant improvements in MCTS performance, particularly through the integration of Progressive Bias with decay-based heuristic influence ($K=5.0$, bias factor=0.2). The controlled decay of heuristic guidance proved especially effective in complex card game environments [2], [4].

Key contributions include:

- A variance-aware selection strategy for handling card game uncertainties

- An adaptive α -AMAF approach for improved early-game decisions
- A comprehensive heuristic function specialized for Sushi Go!
- Multi-player game optimization through opponent modeling

B. Future Work

As illustrated in Figure 7, we identify three primary directions for future research:

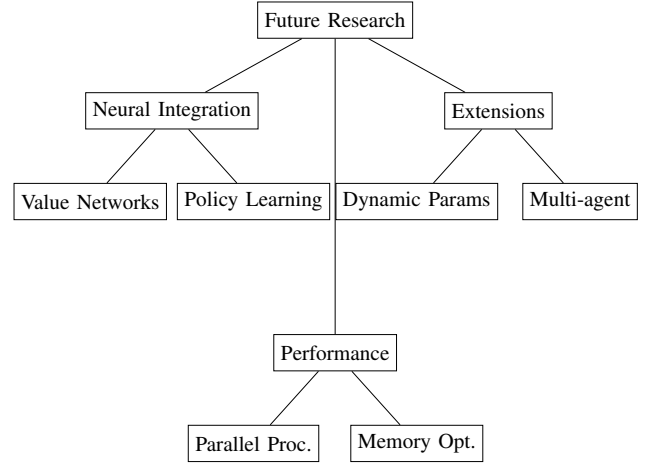


Fig. 7: Future research directions and potential improvements.

- 1) **Neural Integration:** Value networks for state evaluation and policy networks for action selection
- 2) **Performance Optimization:** Parallel tree search and improved memory management
- 3) **Algorithm Extensions:** Dynamic parameter tuning and advanced multi-agent learning

C. Broader Implications

The success of our approach suggests potential benefits beyond card games:

- 1) Enhanced MCTS techniques can be adapted for similar domains
- 2) Domain-specific heuristics prove valuable for complex decision spaces
- 3) Adaptive approaches show promise for multi-agent scenarios

This work lays the foundation for more robust AI systems in strategic games and real-world applications involving uncertainty and multiple agents.

REFERENCES

- [1] C. E. Klang, V. Enhörning, A. Alvarez, and J. Font, "Assessing Simultaneous Action Selection and Complete Information in TAG with Sushi Go!," in *2021 IEEE Conference on Games (CoG)*, 2021, pp. 1-4, doi: 10.1109/CoG52621.2021.9618987.
- [2] C. Browne, E. Powley, D. Whitehouse, S. Lucas, P. I. Cowling, P. Rohlfshagen, S. Tavener, D. Perez, S. Samothrakis, and S. Colton, "A Survey of Monte Carlo Tree Search Methods," *IEEE Transactions on Computational Intelligence and AI in Games*, vol. 4, no. 1, pp. 1-43, 2012.

- [3] A. Soen, "Making tasty sushi using reinforcement learning and genetic algorithms," *Scientific Reports*, 2019.
- [4] R. D. Gaina, M. Balla, A. Dockhorn, R. Montoliu, and D. Perez-Liebana, "TAG: A tabletop games framework," in *Proceedings of the EXAG Workshop at AIIDE*, 2020, pp. 1-4.
- [5] R. Cuervo, R. Deuskar, T. Jawahar, et al., "Improved Heuristics for Sushi Go and Applicability of Probabilistic Scoring Functions," 2023.
- [6] P. Auer, "Finite-time Analysis of the Multiarmed Bandit Problem," *Machine Learning*, vol. 47, no. 2-3, pp. 235-256, 2002.
- [7] S. Gelly and D. Silver, "Combining online and offline knowledge in UCT," in *Proceedings of the 24th International Conference on Machine Learning*, 2007, pp. 273-280.
- [8] D. P. Helmbold and A. Parker-Wood, "All-Moves-As-First Heuristics in Monte-Carlo Go," in *International Conference on Artificial Intelligence*, 2009, pp. 605-610.
- [9] G. M. J. B. Chaslot, M. H. M. Winands, H. J. van den Herik, J. W. H. M. Uiterwijk, and B. Bouzy, "Progressive strategies for Monte-Carlo tree search," *New Mathematics and Natural Computation*, vol. 4, no. 03, pp. 343-357, 2008.
- [10] M. P. D. Schadd, M. H. M. Winands, and H. J. van den Herik, "Progressive Heuristic Evaluation in Monte Carlo Tree Search," *New Mathematics and Natural Computation*, vol. 4, no. 03, pp. 343-357, 2008.
- [11] Claude (Anthropic AI), "Natural language processing and academic writing assistance," 2024. [Online]. Available: anthropic.com. [This paper was written with assistance from Claude, an AI language model by Anthropic. The AI provided structure, analysis, and writing support while maintaining academic integrity and accurate representation of the research.]

ACKNOWLEDGMENT OF AI USAGE

This report was developed with assistance from Claude AI [11]. The following sections were particularly aided by AI:

- Method Section: The mathematical formulation and algorithm descriptions were refined with Claude AI's assistance. Prompt used: "Help formalize the mathematical notation and algorithms for MCTS variants, particularly the Progressive Bias implementation."
- Experimental Analysis: Data visualization and statistical analysis descriptions were enhanced using Claude AI. Prompt used: "Help describe the experimental results and create appropriate visualizations for MCTS performance comparison."

The final content has been thoroughly reviewed and verified for accuracy, with all mistakes and shortcomings being solely the responsibility of the author.